


CVAE-MSA: Fusion-Space Reconstruction Architecture

Training forward pass (text dropped). Deterministic $z=0$ at inference time.



Training: Randomly drop one modality per batch. CVAE encoder uses true h_{target} (teacher forcing).
Inference: $z = 0$ (prior mean). CVAE decoder reconstructs missing h_m from available modalities. No teacher forcing.
Key insight: KL weight (β) dominates all hyperparameters (63% of MissT variance). Optimal β scales with dataset size.