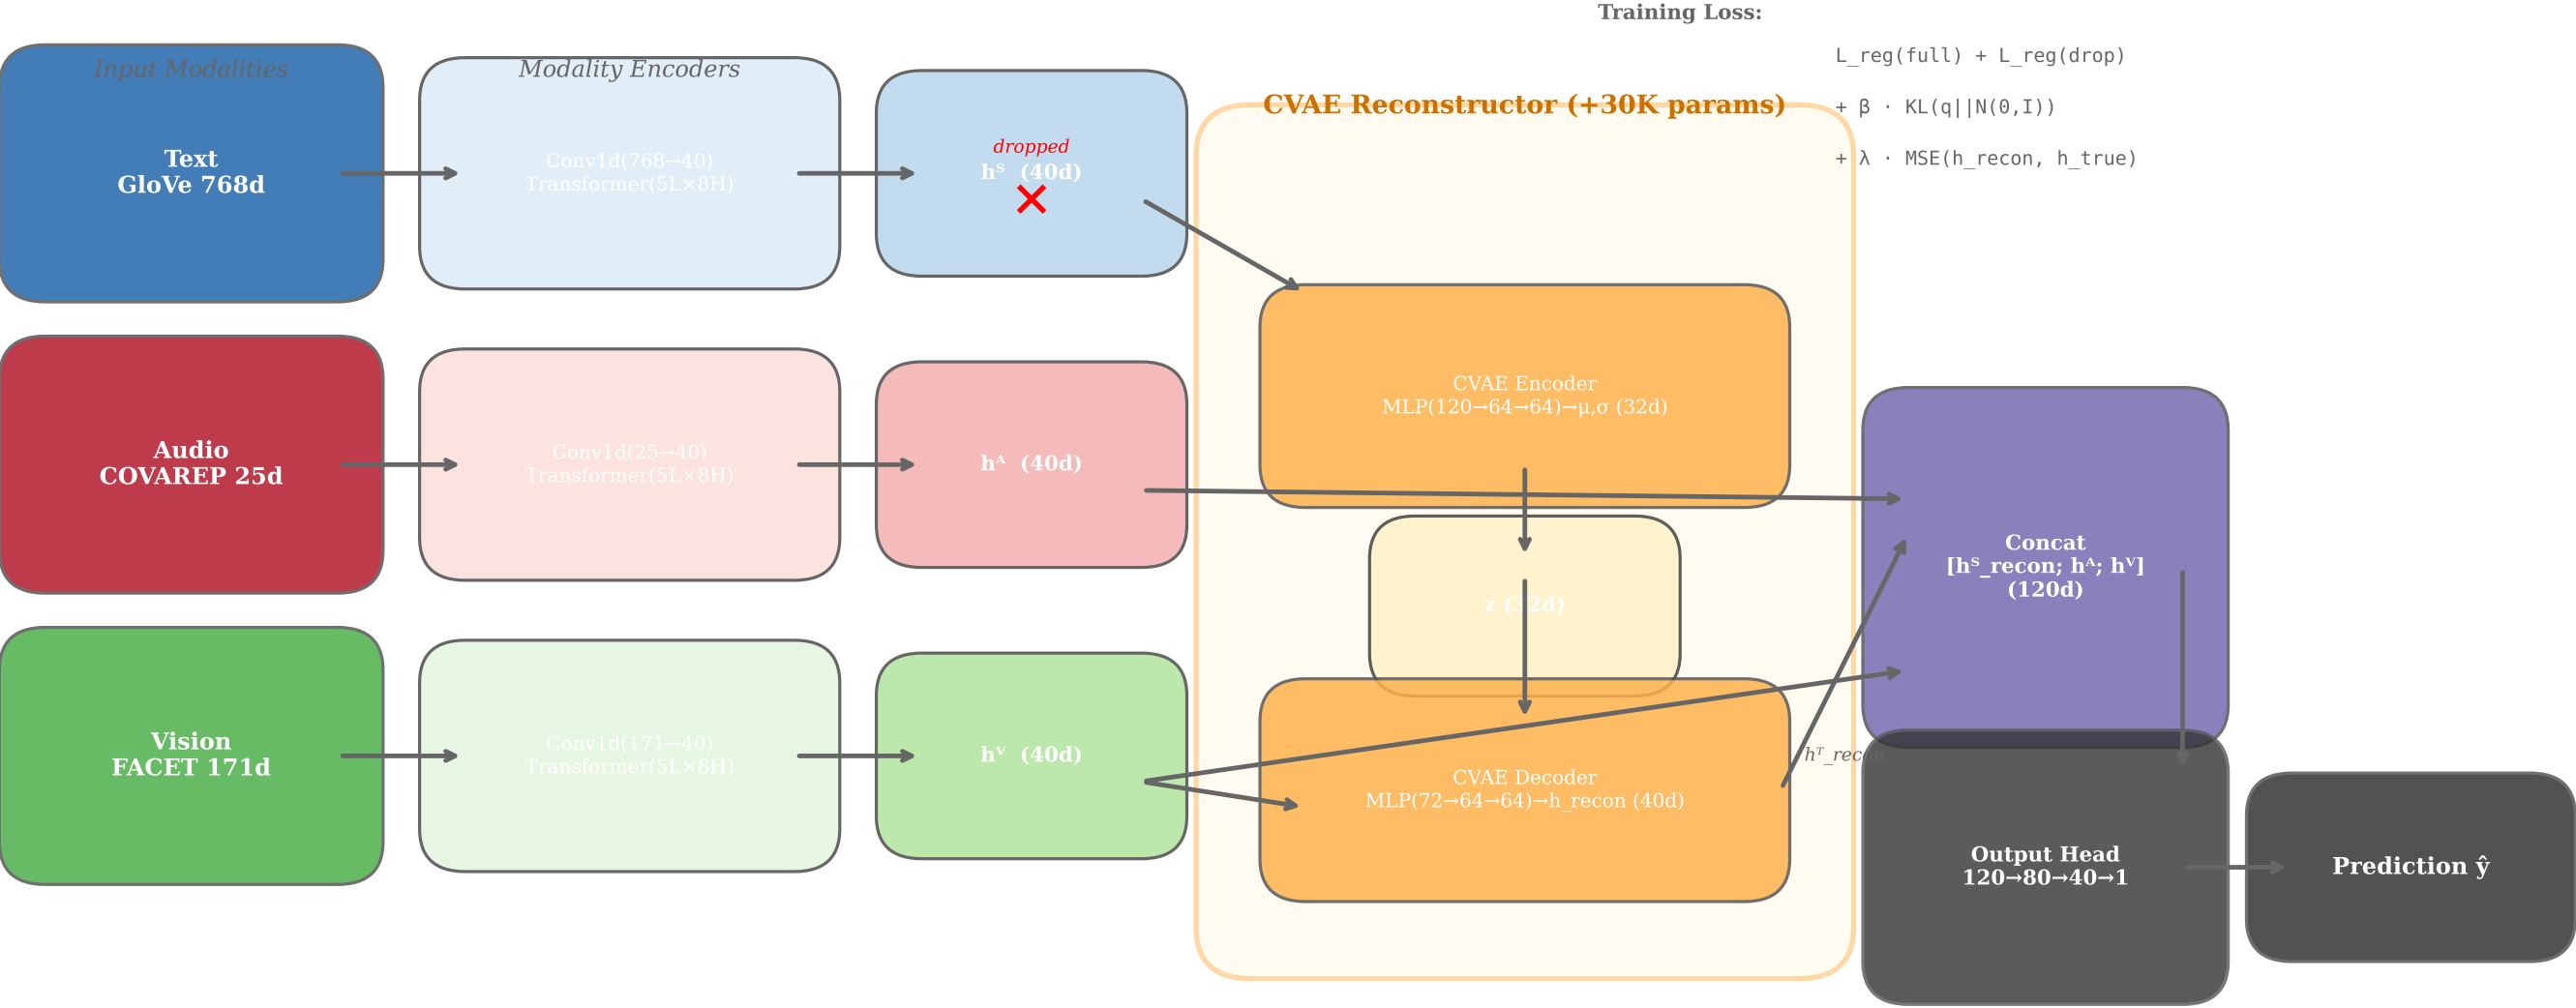


Figure 6: CVAE-MSA Architecture (training with text modality dropped)



Training: randomly drop one modality, use true h_{target} for CVAE supervision

Inference: $z = 0$ (prior mean), CVAE reconstructs missing h from available modalities